

Creating an End-to-end Testing Framework for an Open Source Minecraft Server

A Major Qualifying Project (MQP) Report

Submitted to the Faculty of

WORCESTER POLYTECHNIC INSTITUTE

in partial fulfillment of the requirements

for the Degree of Bachelor of Science in

Computer Science

By:

Alexander Kinley

Matthew Worzala

Project Advisors:

Professor Joshua Cuneo

Date: April 2022

This report represents work of WPI undergraduate students submitted to the faculty as evidence of a degree requirement. WPI routinely publishes these reports on its website without editorial or peer review. For more information about the projects program at WPI, see <http://www.wpi.edu/Academics/Projects>.



WPI

Abstract

We created Canary, an end-to-end testing framework for the open-source Minecraft server Minestom. Canary enables developers to build their tests like they would when testing manually and then save the test structure. Code is used to define the expected outcome of the test. Canary can then load the test structure, and run the test code to determine if the test passed or not. By the end of the project, we have a functional testing framework along with a demonstration of its usage.

Acknowledgements

Begin list of acknowledgements

- First Acknowledgement
- Second Acknowledgement

Contents

1	Introduction	1
2	Background	2
2.1	Automated Software Testing	2
2.2	Unit Testing	3
2.3	End-to-End Testing	3
2.4	Minecraft	4
2.4.1	Custom Minecraft Servers	5
2.5	Minestom	7
2.6	Past Work	7
2.6.1	Mojang GameTest	7
2.6.2	Third Party Tools	9
3	Requirements	9
3.1	Custom Entity or Block Behavior	9
3.2	Custom Server Commands	10
3.3	Customizing Player Interaction	10
3.4	Final Requirements	11
4	Implementation	12
4.1	Test Builder	12
4.1.1	Structure Files	16
4.2	Assertion Engine	18
4.2.1	Soft Pass	19
4.2.2	Assertion Specification	20
4.2.3	Player Interaction Testing	20
4.3	Test Executor	21
4.3.1	Sandbox Versus Headless Mode	21

4.3.2	Test Isolation	22
4.3.3	Test Environment	22
4.3.4	Error Reporting	23
5	Discussion	24
5.1	Validation	24
5.1.1	Canary Demo	25
5.2	Future Work	28
5.2.1	Unified workflow	28
5.2.2	Testing Minestom Internally	29
5.2.3	Behavior-driven Development	29
6	Conclusion	30
References		30
7	Appendix A: Java Classes & Runtime Modification	30
8	Appendix B: Tooling	32

List of Tables

List of Figures

1	Example Minecraft world (https://live.staticflickr.com/8167/7659790494_78158cf095_b.jpg). CC BY-NC 2.0	5
2	Bedwars game in progress, Hypixel, 2021 (https://www.sportskeeda.com/minecraft/top-5-things-players-know-bedwars)	6
3	Custom turkey creature added to Minecraft, by Origin Realms, 2022 (https://originrealms.com/blog/harvesting-update).	6

1 Introduction

Minecraft is a popular video game developed by Mojang Studios and officially released in 2011. The basic premise is that the world is made up of many cube-shaped blocks, and the player has the ability to break and place these blocks. Since its first release, Minecraft has grown into something much larger than this basic premise. One part of this growth is through custom multiplayer servers that allow for entirely new ways to play the game. Traditionally these custom servers have been accomplished by modifying the default multiplayer server. This has been very fruitful, and there are a massive number of servers that take this approach.

Minestom is an open source library that re-implements the basic functionalities of a Minecraft server. The goal of Minestom is to provide a strong foundation on which people can build highly customized Minecraft servers. Minestom is more performant and far easier to extend than the server provided by Mojang because it was built from the ground up with those goals. One of the ways Minestom works towards these goals is by not implementing all of the features provided by the Mojang server. As a consequence, using Minestom makes sense when “removing the features you don’t need takes more time than implementing the ones you do” (Minestom Contributors, personal communication, 2022). People who want to create custom behavior can choose to build it on top of Minestom instead of using other tools that work by modifying the Mojang multiplayer server. Since Minestom is a library that is aimed at people who want to implement custom functionality, providing a robust software development framework is important.

Test-driven development is a software development methodology where the software requirements are first encoded in test cases before they are implemented. This means that you can continually test your software using the test cases and thus have some reassurance that the software is working correctly. This is commonly done using unit tests, which are tests scoped to small parts of the code such as a class or function and have no external dependencies. This is great for making sure that small parts of a larger system are working properly. They are simple to integrate into a Java codebase using libraries such as JUnit. When you want to test the behavior or functionality of an entire system, integration or end-to-end tests are the solution. For instance, this may take the form of automatically clicking through buttons on a form to ensure that it reaches the correct end state, or sending a predetermined set of inputs into a video game and verifying that you end up where you expect.

Minecraft presents a particularly interesting situation with regard to end-to-end testing. The world is made up of different blocks and entities with defined behavior and interactions. Everything is modular and intended to behave the same in many different situations. You can manually test a feature by building some test structure and then taking steps like spawning entities or interacting with blocks to test the feature. You would look for the expected behavior in order to determine if the feature is working as expected. This is a workflow that can be automated. You can build the test structure and save it, then programmatically perform the additional steps such as spawning entities or interacting with blocks. Finally, assertions can be made on the behavior of the situation.

2 Background

Currently, there are no other existing libraries that implement what Canary aims to accomplish for Minestom. That said, software testing and video game development practices are fields with extensive work and research. Software testing is a major aspect of modern software development practices. Because of the benefits it provides, the video game industry has been working to better incorporate automated testing into their workflow. Minecraft has a rich history of modding and customization, including multiplayer servers with features entirely customized and different from the standard Minecraft experience. Minestom is a relatively new project that offers an alternative to previous Minecraft server modding methodologies, allowing for a custom server to be built from the ground up.

2.1 Automated Software Testing

When developing software, it is important to verify that the program works as intended. All software is prone to bugs and instabilities, so ensuring that a program behaves as intended for all expected inputs is a requirement. The most basic way to accomplish this is by manually using the software and confirming that it functions properly. Manual testing is a key part of all software development. It allows the tester to try a wide variety of inputs and operations. The downside of manual testing is that it requires active human involvement, and for large, complex code bases, it is impossible for people to manually test every feature at every stage. Furthermore, complex interdependencies can cause errors to appear in unexpected places.

One solution to this problem is automated testing. Automated software testing encompasses a wide variety of different methods, techniques, and strategies for trying to automatically determine if a piece of software works as expected. Automated tests can be run frequently, ensuring that changes do not introduce unexpected errors. On a high level, automated tests generally involve defining a set of inputs, a series of operations, and an expected output. The thing being tested can be as small as a single function, or as large as an entire distributed system. Automated testing is widely used in the software development industry.

Due to the importance of automated testing, some software development methodologies put testing front and center. One such example is test-driven development (TDD). TDD is a series of steps and guidelines to follow when developing software. The core of TDD involves writing and running small unit tests for a feature before actually writing the code that implements the feature (Bhat & Nagappan, 2006). TDD promises to substantially improve software quality without having a significant impact on overall productivity. The practice is somewhat controversial within the software development community, with varying and inconclusive evidence as to whether it is superior to other agile development practices (Karac & Turhan, 2018). Regardless of the exact benefits of TDD, its existence and supporters demonstrate a general shift in the software development community towards placing higher importance on automated software testing.

Video games are no exception to the need for testing. Current practices in the game development industry involve a large amount of manual testing. This manual testing generally aims to verify that the program behaves correctly in all scenarios. Because of this reliance on manual testing, there have been attempts to automate some or all of that work. As discussed by Politowski et al. (2021) game developers have created many different testing strategies and methodologies for their games. These range from unit tests and lightweight integration tests, to applying machine learning to create smart bots to test their games. One common testing method is scenario-based testing. This method encapsulates a whole family of strategies

involving using a human-defined starting state and a series of inputs or actions (Albaghajati & Ahmed, 2020). The scenario being tested could be playing through a level using pre-recorded inputs, or it can be much smaller in scope such as ensuring that the game renders correctly in certain predefined situations. Scenario-based testing works best when the starting state and actions can be defined once and are expected to continue producing the same output throughout the development lifecycle. For this reason, it is an ideal strategy to apply to games like Minecraft where the behavior of individual blocks or entities will generally remain consistent, even if other parts of the game change.

2.2 Unit Testing

As defined by Huizinga (2007), unit testing is a type of automated testing where tests target the smallest functional unit of a codebase, such as a class or a function. Unit tests are often as simple as giving a known input to a function, and checking that the returned value is as expected. By writing and running unit tests for a code base, you can help verify that future changes do not cause bugs or errors in existing code. Having a large library of unit tests also allows for safer and faster refactoring in the future, since you already have tests to ensure that the code continues to function properly.

A well-tested code base can have hundreds or thousands of individual unit tests. Because of this, it is important that unit tests are simple to write and simple to run. To this end, the Java framework JUnit provides a simple and powerful way to test Java code (Rakshith & Manjunath, 2020). The library handles letting developers define their tests, executing the tests, and reporting on the results. JUnit tests are defined as a method annotated with ‘@Test’. Various assertion functions are provided such as ‘assertEquals’ to check if two values are the same, or ‘assertTrue’ to check if a boolean value is true. When the tests are run, JUnit will execute each of the methods annotated as a test and report if the test passed or failed.

The power of JUnit comes from the fact that it is essentially the default unit testing library for java. Many integrated development environments (IDE) have support for creating, running, and viewing the output of tests in a convenient manner. On top of that, the framework provides additional features such as letting the developer run setup and teardown code before and after each test, or before and after all the tests. The ubiquity and power of JUnit allow it to serve as a model for what an automated testing framework should provide.

This type of automated testing, while vital, is not suitable to be the sole testing strategy. Because the tests focus on small individual sections of the code, they do not test the combined functionality of these units. Unit tests also are not suitable for finding all bugs in the code, as these tests can only show whether the code works properly in all the tested cases. If there is a case that does not get tested, there is no guarantee that it works properly.

2.3 End-to-End Testing

As mentioned, unit testing is not the only type of automated testing and it is not the only type developers should use. Ensuring that all the pieces of an application work together in a correct and robust manner is key to high quality software. End-to-end testing comes closer to mimicking what a developer might do while manually testing their software. By running the entire application and testing the response to user actions or state changes, the functionality of the whole program can be tested. For instance, when a

correct username and password are entered on a login screen, the user should be directed to the homepage. This would be manually tested by navigating to the login screen and entering a known good username and password. Then, if you are shown the homepage you know the feature works. An automated end-to-end test of this feature would work in a similar way, but instead a program might simulate the user inputs, automatically navigating to the login screen and entering a username and password. Finally, checking that it is on the homepage. Like unit testing, end-to-end testing is often facilitated by a framework to allow automation of things such as user inputs, navigating a page, or verifying that the correct result is being shown. In the case of websites, the Selenium browser automation project allows users to programmatically control and get data from a web browser such as Chrome or Firefox (Selenium Project, 2021). This enables developers to create end-to-end tests of their web applications by testing them in the exact environment they will be run in by users: a web browser.

While end-to-end tests are a valuable method for testing an application, larger scope can dramatically increase the complexity and time cost of creating and running them. Firstly, there must be a way to automate the setup, inputs, and actions of the tests. While some platforms, such as websites, are targeted by popular frameworks, uncommon or custom platforms cannot be used in conjunction with existing end-to-end testing tools. This can be made even more complicated by the fact that a system may not be fully deterministic. For example, if a system being tested makes a call to a database, the test must wait for the database to respond before checking the output.

Even with these challenges, end-to-end tests are widely used in industry. End-to-end testing is not a replacement for unit testing, but is instead a method to use in conjunction with unit testing. The ‘test pyramid model’ advocates for viewing automated tests as a pyramid with unit tests at the bottom and high level end-to-end tests at the top (Alger, 2018). The middle of the pyramid consists of integration tests that ensure that certain systems of the program work together, or integrate, correctly. As the pyramid shape implies, there are fewer high level end-to-end tests than low level unit tests, but the success of the high level tests relies on the correctness of the lower level tests. End-to-end testing can be as small as screenshot testing of video games. It can be as large as systems like Netflix’s Simian Army, which is designed to test the robustness of a system by randomly rebooting or shutting down servers or even whole datacenters (Tseitlin, 2013).

Automated software testing is a broad field with a long history. Every day, more and more software gets created. Making software that behaves correctly requires effective and efficient means of testing those behaviors. By looking at the current world of testing both for video games and for the industry at large, we can observe what approaches and methodologies have succeeded. That creates a strong ideological foundation to build new testing frameworks from.

2.4 Minecraft

Originating as a small indie game, Minecraft has become the best-selling game of all time since its release in 2011 (Mojang, 2019). The game has been developed in the Java programming language since its inception; however, it is now split into two versions: Java Edition and Bedrock Edition. Bedrock Edition is developed in C++ and was created after Microsoft acquired Mojang Studios (the developer of Minecraft) for \$2.5 billion in 2014 (Peckham, 2014). Minecraft is a first person game based in a procedural voxel world. Voxels are analogous to 3D pixels and make up uniform cubes in a Minecraft world. These cubes can be

created and broken by the player to create their own structures. In recent years, however, Minecraft has turned into much more than a simple sandbox game.



Figure 1: Example Minecraft world
(https://live.staticflickr.com/8167/7659790494_78158cf095_b.jpg). CC BY-NC 2.0

2.4.1 Custom Minecraft Servers

Minecraft servers have evolved throughout the years, starting from simple survival servers with small extra features to game modes which look and feel like a new game completely. Minecraft itself comes with goals for the player to accomplish in the form of Advancements (a type of achievement), and custom servers often focus on changing the goals of the game. For example, a custom server might remove the resource-gathering aspect of the game to focus on player versus player (PvP) content.

Custom Minecraft servers have become incredibly popular in the past few years. The most well-known of these servers, Hypixel, hosts nearly twenty custom game modes. The server reaches nearly 100,000 concurrent players during peak hours of the day (Minecraft-servers, 2022), making it the most popular independently managed game server of all time (Guinness World Records, 2017). While there is no public data on the company’s income, it has been estimated that Hypixel makes between five million and ten million dollars per year based on employee salary and number of employees. The custom game modes featured on the server include “Bedwars”, a game where players team up to protect their bed while attempting to destroy the other teams’ beds at the same time. This game mode adds new mechanics to Minecraft for gathering resources and fighting opponents. Figure 2.4.1.1 shows a Bedwars game in progress. Some of the modified elements include the scoreboard on the right side and health and team display above each player. Another game mode, “UHC Champions”, spawns players in a limited Minecraft world where they must gather resources to fight each other, but there is a twist. Players can only regain health through hard-to-acquire “Golden Apples”, and the last remaining player wins. This game mode features custom items added to the game such as an axe which strikes lightning when attacking other players (Hypixel Inc., 2022).

While Hypixel features significant modification to the Minecraft server, the gameplay remains similar to the original game. Other popular servers take a more extreme approach to creating new experiences.



Figure 2: Bedwars game in progress, Hypixel, 2021
<https://www.sportskeeda.com/minecraft/top-5-things-players-know-bedwars-minecraft>).

One such example is Origin Realms. Using only server-side modification, Origin Realms creates an experience which feels like a completely different game while still maintaining the same core gameplay. The server adds new blocks, items, creatures, and complete biomes to the game. An example of the server can be seen in Figure 3, which shows a completely custom creature added to the game. These extreme modifications are done through a combination of server software modification and the Minecraft resource pack system. Resource packs allow “players to customize textures, models, music, sounds, languages, texts such as the end poem, splashes, credits, and fonts without any code modification” (*Resource Pack*, 2022). They can be installed on the client, or sent to the client automatically when joining a server. Using resource packs, custom models such as the turkey in Figure 3 can be loaded by the client and moved or animated by the server.



Figure 3: Custom turkey creature added to Minecraft, by Origin Realms, 2022
<https://originrealms.com/blog/harvesting-update>).

2.5 Minestom

Started in July 2019 by TheMode, Minestom aims to reimplement the Minecraft protocol from scratch in Java. Most custom Minecraft servers are created using a modified version of the server published by Mojang, whereas Minestom was started from scratch. This methodology comes with some tradeoffs. When modifying the Mojang server, all of the game features such as entity behavior, block interactions, and much more are already implemented. Minestom, on the other hand, requires the developer to implement all of these features if a server wishes to use them. As such, it makes the most sense to use Minestom when “removing the features you don’t need takes more time than implementing the ones you do” (Minestom Contributors, personal communication, 2022). Using Minestom does not only come with disadvantages, however. Mojang’s Minecraft server has been in development since the game’s release more than 10 years ago, and it contains a huge amount of legacy code which could be optimized significantly. Minestom, on the other hand, was created with a focus on performance, and it does just that. As of April 2021, it can easily handle more than 1,000 players in a basic environment (TheMode, 2021). Since then, internal benchmarks have shown Minestom’s capability to handle more than 5,000 players at once (Minestom Contributors, personal communication, 2022). Additionally, Minestom allows the programmer to work at a packet level to implement features very easily, whereas frameworks such as Spigot require the programmer to work with an unstable API based on deobfuscated code.

2.6 Past Work

Test-driven development has become a popular technique in recent years, and as such this project is not the first to attempt to bring testing into the Minecraft world. Relevant past works include Mojang’s GameTest framework, MockBukkit, McTester, and a prototype internal Minestom testing tool. Unfortunately, none of these solutions meets the needs of high level testing in the Minestom ecosystem.

2.6.1 Mojang GameTest

Mojang itself has created two similarly named tools for testing both the Java and Bedrock Editions of Minecraft. Both tools are named GameTest, so they will be referred to as Java GameTest and Bedrock GameTest for the purpose of this report. Java GameTest is completely internal, whereas Bedrock GameTest was made available to the public in June, 2021 (Ammerlaan, 2021). Both frameworks operate on the premise that a developer can create a structure inside the game, and write some code to operate on the structure as well as make assertions. Listing 2.6.1.1 shows a sample test from the Bedrock GameTest framework. As seen on lines 14-16, a test is registered with a structure to contain the test, and a function to set up the initial state and make assertions. The example test requires that the victim—a fox in this case—is in the test area when it starts (Line 8) and not when it ends (Lines 9-11).

```
“‘js
// Listing 2.6.1.1
// Adapted from https://docs.microsoft.com/en-us/minecraft/creator/documents/gametestbuildyourfirstgametest
function simpleMobTest(test) {
```

```

const attackerId = "fox";
const victimId = "chicken";
test.spawn(attackerId, new BlockLocation(5, 2, 5));
test.spawn(victimId, new BlockLocation(2, 2, 2));
test.assertEntityPresentInArea(victimId, true);
test.succeedWhen(() => {
    test.assertEntityPresentInArea(victimId, false);
});
};

GameTest.register("StarterTests", "simpleMobTest", simpleMobTest)
.maxTicks(410)
.structureName("startertests:mediumglass");
```

```

The Java version of GameTest behaves in a similar manner; however, it uses Java for test functions instead of JavaScript. According to Henrik Knilberg (2020), a member of the Java gameplay team at Mojang, the test framework has a few central goals. Tests must be easy to read, write, and run, they must run visibly in an actual Minecraft world as well as on a build server, and they must be quiet on success but verbose on failure. Figure 2.6.1.2 shows some examples of what the tests look like. Individual test cases are inside the white bounding boxes, and a concise output is sent to the player when tests pass. This model serves as a great inspiration for creating and running tests inside the game, but it has significant flaws for widespread use. It is not made for Minestom and therefore makes a number of assumptions about Mojang features being implemented. Furthermore, the framework is closed source, so modifications are challenging, ineffective in a different environment, and cannot be distributed.



Figure 2.6.1.2: Example tests from the Mojang (Java) GameTest framework

## 2.6.2 Third Party Tools

In the category of modifications to the Mojang server, there have been two major attempts at automated testing: MockBukkit and McTester. Both take a programmatic approach to testing and target different use cases. MockBukkit provides the user with an API for creating mocked resources such as players, worlds, and Bukkit plugins (Schaetzen, 2022). This approach allows for a variety of testing methods; however, since the tests are created programmatically, even simple tests are time consuming to write. McTester works by running a modified Minecraft client which allows the test code to submit commands to the client. The server can then make assertions about the result of the client actions server side. For example, a test could include a client right-clicking a Command Block and asserting that the GUI was only opened if the player had appropriate permission (Hill, 2022). This method of end-to-end testing is effective for testing a server implementation, such as testing Minestom internally. Unfortunately, it introduces unnecessary overhead for testing applications built on top of Minestom. Applications generally do not interact with the client directly, and the actual server internals are expected to be tested and working independently. Finally, Minestom has seen its own attempt to create a testing framework internally. This framework involved a mocked client connection to the server, from which a test could submit packets and make assertions on the received packets. Similar to McTester, this can be effective for testing internal features; however, it becomes cumbersome to work on a packet level when testing user created features.

## 3 Requirements

As previously discussed, this project aims to fill a void in the Minestom ecosystem for testing solutions. We hope to create a library to be the de facto standard for end-to-end testing in the Minestom ecosystem. In order to meet this goal, we need to consider two main factors: ease of use and coverage of major use cases. Ease of use means that developers using Canary should not feel like they have to fight with the library in order to accomplish what they want. The library should be internally consistent and predictable, while handling the needs of developers. To cover major use cases, Canary should be applicable to the scenarios where developers would want to write end-to-end tests.

To figure out what standard usage might look like, we looked towards the greater Minecraft server modding community to see what sort of features or functionality people have implemented within a Minecraft server. Although the functionality and features that people add to Minecraft servers are broad and diverse, we can break them down into a few general categories based on what they require from a testing library. Depending on how a feature affects the world, it may require different approaches for testing. We will look at three broad categories of functionality that can be added to a Minecraft server.

### 3.1 Custom Entity or Block Behavior

This category refers to most world behavior that is not linked to the player. This includes entity AI, entity properties, block properties, how blocks interact with other blocks, and more. Generally these features only affect a small area around the block or entity, and they happen based on the state of the world

and not the actions of players in the world. Manually testing these types of features would generally involve building some sort of setup that will cause the desired behavior to occur and then watching for the behavior to happen as expected. For example, to test that minecarts roll down hills correctly, the developer would first build a track down a hill and then place a minecart on that track. Then they would watch for the minecart to end up at the bottom of the track. Canary allows developers to define the setup (structure) for a test, as well as the expected outcome, and then execute the test on its own. The required functionality is as follows:

- Define the starting state of a test, including the blocks, block properties, and entities.
- Check if expected outcomes happened.
- Run the test until it passes or fail if it has not passed during a set lifetime.

Determining if the expected outcome happened, in general, is a challenging problem. There are a wide variety of things that someone might look for to determine if a test ran correctly. Examples include looking for an entity to get to a location, looking for a block to be in a certain state, making sure some condition did not happen, or checking for things to happen in a certain order. Section 4.2 discusses assertions more in-depth.

This category of feature benefits the most from end-to-end testing. Because most of the behaviors only cause changes in a small area, tests can be run simultaneously without having to reset the entire world between tests. On top of the functionality requirements, the process of making, running, and debugging these tests should be simple and straightforward to align with the overall goal of ease of use.

## 3.2 Custom Server Commands

Server commands are a way to use player messages like a command line or terminal. Custom commands are a common way to implement features like switching between worlds, interacting with server mods, or configuring settings. Server commands, unlike entity and block behavior, frequently affect things far beyond the player, including internal server state. To test server commands we should be able to do the following:

- Execute a command from a simulated player with a known starting state.
- Check for the expected behavior of the command.

Server commands present a particular challenge because their behavior is not bound in any way. A server command can do anything from changing a block near the player, to teleporting the player, to affecting other players, worlds, or even the entire server. This broad scope makes it challenging to be able to effectively test all varieties of server commands, because creating a known starting state could require restarting the entire server between every test, which would seriously impact the ability to quickly run tests.

## 3.3 Customizing Player Interaction

If a Minecraft server is aiming to alter the player experience, they will likely have to change how the player interacts with different aspects of the world. These changes could be anything from teleporting

the player when they click on an item in their inventory to altering the world based on how many enemies the player has killed. Custom player interactions, like server commands, can affect all aspects of the world and the players in it. Enabling developers to test player interactions requires the following:

- Simulate a player from a known starting state in a known context interacting with the world in some particular way.
- Check for the expected behavior of the interaction.

Player interaction often causes changes that can be challenging to test. The changes may interfere with other tests, or require careful setup and teardown to ensure the test succeeds when run multiple times. That being said, some types of player interaction, such as combat or how players interact with blocks, may be easier and more straightforward to test.

### 3.4 Final Requirements

There are many challenges to creating a test framework that covers all expected use cases. Features such as server commands and custom player interaction can potentially change the server in ways that affect the outcome of other tests. Being able to test every possible feature in a consistent and reliable manner would require restarting the server between every test. This would substantially slow down running tests as well as make it harder to see and debug failing tests, forcing the library to be much less useful overall. For this reason, we chose to focus mainly on testing custom entity and block behavior features which can be tested in a confined area. These features can require elaborate setups to fully test, and stand to benefit the most from automated end-to-end testing. That being said, we have not entirely ignored other feature types, and they could be the focus of future work. From this, we can create a list of the high level requirements of Canary. We have split the requirements into two categories, based on our original goals of ease of use and full feature coverage.

#### Ease of Use

- Canary must be easy for developers to use, including creating, running, and debugging tests.
- Allow developers to create tests in a way similar to how they would when manually testing.
- The code API used when writing code for tests should be understandable and abstract away common operations
- Silent on success, verbose on failure.
- Tests should be fast to run on a local machine
- Tests should run on Continuous Integration/Continuous Development (CI/CD) servers.
- Integrate nicely with version control.

#### Feature Coverage

- Canary must be applicable to all major use cases; in particular, custom block and entity behavior should be able to be tested by Canary.
- Developers should be able to express complex scenarios in their tests.
- Cover most common usage and allow for easy developer extension.

## 4 Implementation

With the high level requirements determined, they must be translated into specific requirements to be implemented in software. The approach takes inspiration from the Mojang Java GameTest framework. The key concept from this work is the combination of a simple structure and code to create assertions for the test. This approach allows for the creation of tests to be done largely inside Minecraft, the same way that someone would when manually testing. Additionally, it reduces the amount of code required for each test, further improving ease of use. The implementation of this system can be split into three broad categories: the test builder, assertion engine, and test executor.

### 4.1 Test Builder

As discussed, Canary tests have a test structure associated with them. The structure is not defined by code, it is data that is loaded from the filesystem into the world when a test is being run. To make Canary easy to use for developers, we allow them to create these structures in the same way that they would when manually testing: by building them in Minecraft. The test builder is the subsystem of Canary responsible for letting users create and manipulate these structures. It allows users to build structures in game that will then be saved in a form that can be ingested when running tests.

The test builder allows for the creation of all the necessary starting conditions for tests. It is easy to use and has convenience features that users frequently need when creating end-to-end tests, such as basing the structure for one test off of the structure of another. The created structures allow users to edit the properties of the blocks, either to change how they behave or to mark them so their position can be referenced from the code associated with the test (mark a block with a name that can be used to easily get the position of the block). Because the test builder runs on the server, it has multiplayer support. This means that multiple players can work on a single test structure at once and that different players can work on different structures simultaneously.

Once a user has built a test structure, it is saved in a format that can be loaded into the world when running tests. For git compatibility, structure files are text-based so that minor changes to the structure will not cause git to see the whole file as changed.

When using the test builder, a player is brought into a separate empty world containing just the structure they are creating or editing (Figure 4.1.1). The advantage of this approach is that it becomes simple to track all of the blocks that are part of the structure, and there is no risk of other parts of the world interfering with creating the structure.




Figure 4.1.1: Example of a structure in the test builder

The way that the player controls the test builder—including making a new structure or leaving the test builder when they are done creating their structure—is through server commands. Server commands are a feature of Minecraft that are similar to a command line. Using the in-game chat, messages starting with a forward slash are considered commands. Minecraft and Minestom have powerful support for commands, including named arguments and autocomplete. As seen in Figure 4.1.2, Canary can suggest possible valid inputs when they are known.




Figure 4.1.2: Autocompletion of existing test structure when using the duplicate command.

The test builder commands available to the player change depending on whether or not they are currently in the test builder. When the player is not in the test builder, the commands available to them are as follows:

```
'/test builder new STRUCTURE_NAME'
```

Places the player in a test builder world with the default starting structure. When they save the structure, the file will be named STRUCTURE\_NAME.json

```
'/test builder duplicate STRUCTURE_NAME NEW_STRUCTURE_NAME'
```

Places the player in a test builder world with structure from STRUCTURE\_NAME.json. When they save the structure, the file will be named NEW\_STRUCTURE\_NAME.json

```
'/test builder join CURRENT_STRUCTURE_NAME'
```

When there is a player currently building a test, this command is how you join them in the test builder.

When the player is in the test builder, the commands available to them are as follows:

```
'/test builder done'
```

Saves the current structure and returns the player to wherever they were when they entered the test builder.

While in the test builder, players can freely place and remove blocks. The creative mode of Minecraft gives players the ability to place any block in the game by choosing it from a list. To help users visualize the structure that they are building, a white bounding box is placed around all of the blocks that are a part of the structure. This box is updated as players place and remove blocks to always bound the entire structure. This is accomplished using a Minecraft structure block. A structure block is normally used during world generation in vanilla Minecraft, but for our purposes the white bounding box is all we need (Figure 4.1.3). A structure block is placed below the build area, and its parameters are updated as necessary to change the size and position of the bounding box that it creates.




Figure 4.1.3: The bounding box tracks the size of the structure

The other important functionalities within the test builder are editing the properties of blocks and placing markers. Blocks have data associated with them beyond just what type of block they are. This includes information like the orientation of the block, what Java class is responsible for handling the block's behavior, and many other properties. The fields cannot normally be set in Minecraft; they generally are set by the game itself. Markers are a concept introduced by Canary as a way for a user to give a name to a position within the test structure. The marker name can be used to get the location of the marker in the code for a test. Markers allow for a more ergonomic alternative to putting raw coordinates in a test.

These two functionalities were implemented using the same general principle, which is to give the player items with special behavior when used to click on an item. While Minecraft does not have a built-in method to prompt the user for typed input, there are multiple different methods that can be implemented to approximate the desired behavior. Two such examples are using the in-game chat and making use of the anvil inventory screen, which has a box for text entry. Using these, we can prompt the user to enter information when they click on a block with one of the special test builder items. For example, with the set marker item, after clicking on the block they want to mark, the user will be given an anvil prompt (Figure 4.1.4) where they can enter the desired marker name. Then they can click on an item (green glass) used as the confirm button or another item (red glass) used for canceling. This will then create a marker with that name at that location in the test structure. A similar process can be used for the various types of block data.




Figure 4.1.4: Example of the anvil prompt for markers

Behind the scenes, all of these special items pass the position of the block that was clicked on—along with the data the user entered—to commands that are part of the test builder. These commands can also be used instead of items if desired by a user. The two commands currently implemented by Canary are for setting a marker and setting the block handler property of a block. They are defined as follows:

`'/test builder edit marker X Y Z MARKER_NAME'`

Creates a marker at position (X, Y, Z) with the given name.

`'/test builder edit handler X Y Z HANDLER_NAME'`

Sets the block handler for the block at position (X, Y, Z) to be the one defined by the given name.

#### 4.1.1 Structure Files

As previously mentioned, the test structures created using the test builder are saved using a text format, specifically a JavaScript Object Notation (JSON) file. This format was chosen because it is both human and machine readable, and is reasonably compact. The fields within the file use a format custom to Canary that was created to be easy to create and parse while trying to minimize the size of test structure files. Large software projects can have hundreds or even thousands of tests, and an overly verbose format could cause test structures to use up an unnecessary amount of disk space. Listing TB.1 shows a sample test structure file.

```
“‘json
// Listing TB.1
{
 "id": "my-test-world",
```

```

 "size": [
 16,
 16,
 16
],
 "markers": {
 "diamond_block": 123
 },
 "blockmap": [
 "minecraft:stone",
 "minecraft:cobblestone_stairs[facing=north]",
 {
 "block": "minecraft:stone_stairs[facing=south,waterlogged=true]",
 "handler": "example:my_block_handler",
 "data": {"name1:123,name2:\"sometext1\",name3:{subname1:456,subname2:\"sometext2\"}}"
 }
],
 "blocks": "0,256;1,16;0,240;-1,3584"
}
```

```

The fields of the JSON file in Listing TB.1 are specified as such.

id

Some unique id to be referenced from the test code.

size

The x, y, and z size of the structure, with a max of 48, 48, 48.

markers

An object where each key-value pair represents the name and position index of the marker. The position indexing uses the same scheme as described for the block's value.

blockmap

An array where each element represents a block which is used in the structure; a block id of -1 represents a block of air with no additional properties. A block can have a type, tags, a handler, and additional data. A block object has a block, handler, and data field. The handler and data fields correspond directly to the handler and additional data of the block they represent. The block field represents both the block type and the block tags as a single string of format 'block_type[tag1=tag1_value,tag2=tag2_value,...]'

(if there are no tags, the square brackets are omitted). If a block has a special handler or data, it must be defined as using a block object. If a block just has a type and tags, it can just be stored using a string equal to what the block string of its object would be.

blocks

A run length encoding of the blocks for the test. The semicolons separate each block definition. A block definition starts with a block id reference the blockmap array, followed by a number of blocks. Since the structure is three dimensional, but our block definitions are linear, we need a way to convert between the two ie. find the x, y, and z coordinates of the n-th block.

Firstly, we define the position (0, 0, 0) as the first position in the structure. From one position, the next position can be gotten by incrementing the x coordinate. If the new x coordinate is greater than the size of the structure, you set the x to be zero and increment the z coordinate. Likewise if the z coordinate is greater than the size of the structure, you set the z to be zero and increment the y coordinate. The order of positions will generally look like (0, 0, 0), (1, 0, 0), (2, 0, 0), ..., (0, 0, 1), (1, 0, 1), ..., (0, 1, 0), (1, 1, 0), ..., assuming that the structure is large enough to encapsulate all of these positions.

The run length encoding used in the blocks field is the primary way that this format helps to reduce file sizes. Most test structures will likely contain only a handful of different block types, with a lot of repetition. The run length encoding—along with the block map array—means that much of that repetition can be stored succinctly. This format is not the most optimal way to compress this information, but we have found it does well enough in most use cases, while being easy to output and parse.

4.2 Assertion Engine

Assertions and their associated Application Programming Interface (API) are an integral part of Canary and one of the most user-facing elements. They are how the user encodes all expected functionality for the tested features of their software. In order to meet the project goals, the API must be highly flexible while remaining simple for obvious tasks. Assertions in Canary do not happen in a deterministic manner, and they may happen at an unknown time from the start of the test. This is because the Minestom server does not run at an exact rate, so there may be slight differences in behavior between executions. This creates a significant challenge for writing assertions which remain simple to understand. The solution adopted by Canary is to allow the user to define small programs which dictate the expected state of the test environment throughout its execution time. The program can be conceptually split into two distinct segments, roughly modeled after a typical interpreted programming language: assertions themselves (syntax), and the backing assertion nodes (AST). The syntax of a typical assertion in Canary can be found in Listing 4.2.1.

```
``java
// Listing 4.2.1
expect(myZombie).toBeAt(3, 1, 3).and().toHaveHealth(20.0);
``
```

The syntax is legible like plain English, e.g. "I expect that my zombie will be at <3,1,3> and have 20 health". In this report, an assertion may be referenced using the following abstract representation: EXPECT myZombie.pos=<3,1,3> AND myZombie.health=20.0, or simply pos=<3,1,3> AND health=20.0

for the general case. When the test containing Listing 4.2.1 is executed, the calls generate a list of steps to the assertion. This step is analogous to tokenizing an input source in a programming language. The generated list for this case is shown in Listing 4.2.2.

```
```
// Listing 4.2.2
SUBJECT = myZombie
STEPS = [
 subject.pos == <3, 1, 3>,
 AND,
 subject.health == 20.0,
]
```

```

From the step list, the engine generates a tree of nodes, each of which has the ability to report whether it passed or failed depending on the condition (eg `subject.pos == <3,1,3>`) or the children. For example, the AND node has the following simple logic:

1. If any child is a FAIL, return FAIL
2. Return PASS

This system is analogous to an Abstract Syntax Tree (AST) in an interpreter, and it allows the test executor to simply check the root node to determine the result of the assertion. The nodes recursively process until all branches have been evaluated or an early exit has been reached. The test executor acts on the result of the assertion tree, as described in Section 4.5.

4.2.1 Soft Pass

The assertion engine as described above is effective; however, it is not perfect. A common assertion tool is the NOT statement. This allows the user to invert the statement directly following. For example, `NOT health==20.0` says that it will fail if the subject's health is equal to 20. By defining what causes an assertion to fail, it becomes difficult to determine when it has passed. Assertions are executed every server tick until they return a PASS (further described in Section 4.5). If, for example, the subject has a starting health of 10, then on the first frame this assertion will return PASS and never be tested again. The solution to this problem is to introduce a third state to the equation: `SOFT_PASS`. A soft pass indicates that the condition is currently passing; however, it could fail in the future, so it must be tested again. This system allows for statements which cannot produce a final output. Another use case of this system is the `ALWAYS` statement. It says that the following condition *must* be true for the duration of the test. The `ALWAYS` node can never return a PASS, because the condition could still change in the next tick, so instead it returns `SOFT_PASS`.

4.2.2 Assertion Specification

In order to create the English-like syntax Canary uses for assertions, a somewhat complicated inheritance hierarchy is used. Each call to expect returns an assertion class for the specific object passed in, such as an EntityAssertion if a Minestom Entity is passed in. To reduce code duplication, assertion classes themselves create an inheritance hierarchy, meaning it is possible to call ObjectAssertion.toBe(Object) on an EntityAssertion. This works fine in isolation; however, it falls apart when considering chained assertions such as an AND statement. Each assertion function returns itself so that more statements can be chained together, but control statements are defined on the base assertion class. Consider the code snippet in Listing 4.2.2.1; the and() call returns a BaseAssertion instance, which does not have toBeOnFire(). Java does not have a self type like in some languages, so we must use recursive generics to accomplish a similar effect. Recursive generics are generic types which reference themselves. The (abbreviated) definitions for BaseAssertion and EntityAssertion can be seen in Listing 4.2.2.2. BaseAssertion requires a generic parameter of some type inheriting from BaseAssertion. In EntityAssertion, we fill this requirement. As a result, the return type of not() inside EntityAssertion is still EntityAssertion. While this solution does accomplish the desired goals, it is an unsafe and generally discouraged practice because there is no strong validation on what is being set as the generic parameter. This system is internal to Canary, so these contracts can still be met to create a usable syntax and stability.

```
```
``java
// Listing 4.2.2.1
expect(myEntity).toBeAt(3, 1, 3).and().toBeOnFire()
```

```
``java
// Listing 4.2.2.2
public class BaseAssertion<This extends BaseAssertion<This> {
 public This not() { ... }
}
public class EntityAssertion extends BaseAssertion<EntityAssertion> { ... }
````
```

4.2.3 Player Interaction Testing

Most custom Minecraft servers revolve around player interaction rather than events which happen passively. Creating fake player events for tests can be somewhat limiting, given the complex back-and-forth between the client and server. Manually mimicking client behavior in an accurate manner would be extremely difficult and error-prone, along with requiring a lot of effort. Instead, Canary proposes recording player interactions on a live server and playing them back with a simulated client during a test case. Instead of spawning an entity, for example, a tester could spawn a player with a packet stream that was "programmed" in-game inside the test structure itself. The Canary implementation in this report contains a prototype of this system via a recording mechanism. There is, however, more to be done for a complete system.

The recording infrastructure in Canary uses a simple binary file format (named Canary Packet Record, or CPR). The file format is a short header followed by a list of the binary protocol data received by the server. The header contains the file format version, the Minecraft protocol version, and the starting state of the Player. The Minecraft protocol version must be stored because packet IDs and format change between versions, meaning the file must be updated. Each entry in the file contains the time (since start of recording, in milliseconds) that the packet occurred, and the packet data itself. A detailed description of the file format can be found on GitHub. Once loaded, a fake player is created which submits the packets to the server according to the time they occurred in the CPR file.

The prototype implemented in Canary is far from complete, however. Playing back a recorded interaction requires the tester to manually load the recording and create the playback player. This system should be a simple test environment request, which creates, plays, and cleans up the playback player. Additionally, there are cases, such as player versus player (PvP) features, where it would be beneficial to have multiple players recorded in the same session. For this case, the tester should be able to play one or more recordings while creating a new one so that they are able to interact with the existing recording. For example, hitting it with a sword to ensure that damage is dealt in player versus player combat. Another issue with the packet recording mechanism is the initial Player state.

4.3 Test Executor

The Canary test executor is responsible for managing the testing environment and complying with the JUnit test engine specification. When the engine is invoked, it first discovers every potential test case (method annotated with `@InWorldTest`). Each test is loaded into the appropriate Java ClassLoader, and a Minestom instance is created. An "instance" in Minestom is analogous to a Minecraft world and allows the test to be completely isolated from all others. When the tests are executed, each test structure is placed into its respective instance, and the assertion engine parses the step lists (Listing 4.2.2) into assertion trees. Finally, the test instances are ticked until it receives a definitive result. The stop conditions for a running test are as follows:

1. Every assertion has returned a PASS value (pass)
2. Every remaining assertion returns a SOFT_PASS value (pass)
3. The timeout limit has been reached (fail)

The second condition is due to the behavior of SOFT_PASS. An assertion such as `ALWAYS` cannot ever know that it has definitively passed, so it exclusively returns SOFT_PASS or FAIL. When a test times out, however, it means that there are no more steps so a SOFT_PASS is equivalent to a PASS. Because Canary tests run over a long period of time and state changes throughout, FAIL does not indicate a definitive failure. Instead, it indicates that the test has *not yet passed* but may in the future. If the time runs out and there are still failing assertions, then the test is a failure.

4.3.1 Sandbox Versus Headless Mode

Canary has a strong focus on ease of use and visual feedback through the sandbox mode of the test engine. Writing tests inside the environment they are testing provides the user a huge benefit, including real-

time feedback. The user can enter the game and watch their tests running in real time, making it extremely simple to determine what is wrong. Canary facilitates this using the "sandbox instance", a Minestom instance containing every loaded test. This world allows the user to visibly see which tests are passing or failing, and the user can fly around in the virtual space to view them. Visual feedback for the developer, however, is not the only responsibility of a test engine. Test engines must be able to execute remotely in a Continuous Integration/Continuous Development (CI/CD) pipeline. Canary refers to this as "headless" execution mode, and it differs from the sandbox mode in a few ways. The running test server is not available to join, and the debug information is not loaded, including commands and the sandbox instance. To ensure that all tests remain isolated and predictable, each test is still loaded into its own instance.

4.3.2 Test Isolation

It is critical to create the same exact environment for every execution of a test both locally on multiple distinct machines or remotely as a part of a CI/CD pipeline. This task is non-trivial in the Minestom environment, since a test may access and use arbitrary server data or interact with arbitrary blocks around the test environment. The most obvious solution is to use a unique Minestom server for each test which solves the global state and nearby interaction issues in a simple and intuitive way. Unfortunately, this solution has significant drawbacks in terms of speed and usability. Minestom makes use of a global (static in Java) state for running server information, which means that you cannot easily have two Minestom servers running at the same time. It is possible to start a Java Virtual Machine (JVM) subprocess for each test or load each test into its own ClassLoader. Both results incur a large initialization overhead and memory increase for each individual test because every class must be reloaded each time. Regarding usability, separating each test into its own JVM or ClassLoader means that it is no longer possible to create the aggregate sandbox instance, destroying the ease of creating tests. Instead, Canary handles isolation by placing each individual test (and associated structure) into its own unique Minestom instance. This allows the sandbox instance to remain by forwarding actions from each sub instance into the sandbox. This solution, however, does not impose strict isolation for global server state. In the end, this compromise was worth the benefit provided by the sandbox instance. Without strict isolation, users are required to write code which can be tested in isolation. As a result, the code quality when using Canary should increase.

4.3.3 Test Environment

To achieve the goal of simplicity from a user perspective, Canary provides a utility for various tasks inside the test environment. This class, `TestEnvironment`, provides safe access to the Minestom instance for common tasks such as spawning entities. In sandbox mode, tests may be executed multiple times, so they must be cleaned up properly to avoid creating false positives or negatives from leftover changes to the environment. To accommodate this, the test environment provides a tracking interface. A user may register a trackable object, which will have a cleanup method called whenever the test needs to be reset. Internally, this mechanism is used to remove any entities which are spawned during the course of the test execution. The test environment also provides access to the markers defined inside test structures, without the user directly interacting with the structure file itself.

4.3.4 Error Reporting

In any testing framework, good error reporting is a vital element to success and ease of use. Canary is no exception, but this presents a complex situation due to the long-running nature of the tests. In most testing frameworks, errors are presented as the difference between the expected and actual result. Canary assertions can be complex and not necessarily dependent on any single tick of the test. For example, consider the assertion pos=<3,1,3>. If the subject never reaches the position <3,1,3> before the time runs out, then presenting the user with the final position is not necessarily helpful. For cases such as this, Canary records a log file every frame containing the output of each assertion. This data can be hard to parse, so it should not be considered as the first resort. Headless and sandbox modes provide different environments for errors to be presented, and these are handled individually in Canary.

In headless mode, the output is going to a terminal, meaning the result can be somewhat long and must encode all relevant information into text. In this case, Canary logs the failed assertion tree for the failing tick into the terminal, along with the complete log. An abbreviated failure in headless mode can be seen in Figure 4.3.4.1. In a CI/CD pipeline, users typically do not have access to files produced during execution, so writing a log file is not helpful. Additionally, headless mode does not give any visual output of the test, so the log is the only available information. As such, it is recommended that users instead download the test and execute it locally for an interactive visual output of the running test.

```
““  
// Figure 4.3.4.1  
Testing complete.  
1 test failed.  
Error: com.mattworzala.canary.demo.TestEntityTest.testWalkToDiamondBlock - e_e9f02b@<2, 1,  
3>  
e_e9f02b did not reach <2, 1, 3>  
Complete Log  
e_e9f02b@<2, 1, 3>  
Step 0 : entity1 was at <2, 1, 1>  
Step 1 : entity1 was at <2, 1, 1>  
Step 2 : entity1 was at <2, 1, 1>  
Step 3 : entity1 was at <2, 1, 1>  
...  
““
```

In sandbox mode, the user can see and interact with the test output. This allows the user to debug the test in the same way as someone might do when manually testing a feature. In addition, sandbox mode provides only the last tick failure to the player, deferring to the longer log file for the complete output of the test. A sample failure report generated by the sandbox mode can be seen in Figure 4.3.4.2. The message contains a clickable element ([LOG]) to open the log file locally, and to teleport to the test in the sandbox

to view the result visually. The test name can be clicked to teleport to the test in the sandbox instance. In addition to a more concise output, the sandbox instance indicates test results using a Minecraft beacon. As seen in Figure 4.3.4.3, this is visible from far away allowing the user to quickly identify failing tests in the sandbox.



Figure 4.3.4.2: Report generated in game when an assertion fails.




Figure 4.3.4.3: Red and green beacons indicating success and failure of a test.

5 Discussion

We set out to create an end-to-end testing framework for Minestom capable of becoming the de-facto standard. Towards this goal, we addressed many of the challenges associated with end-to-end testing frameworks, and leveraged the narrow scope of exclusively testing Minestom applications to create a unique and novel system. To be able to claim that Canary lives up to the stated goals, we must be able to assess it in its current state. Beyond that, we must also acknowledge the places where Canary currently falls short and how Canary might be expanded in the future to address other goals.

5.1 Validation

A library is only as valuable as the benefit it gives users. So when developing a software library, it is valuable to gather feedback and opinions from potential users on their thoughts and feelings. On top of

that, a testing library like Canary should ideally be battle-tested and hardened against all varieties of edge cases. That being said, the potential user base for Canary is currently quite small. While Minecraft is a massively popular game, Minestom is a relatively young and new library, and it will likely be a while until it is a truly viable alternative for developers creating custom Minecraft servers. That being said, there are people using and working with Minestom.

Informal conversations with developers using Minestom have indicated that most do not employ automated testing of any kind. While we believe that this is all the more reason for Canary to exist, it means that the validation of the functionality and usability of Canary is limited to a relatively small demo written by the same people who created Canary. Creating this demo let us put ourselves in the shoes of someone using the library.

5.1.1 Canary Demo

Canary was designed and built for the goal of testing and entity and block behaviors and interactions. In order to demonstrate this, we chose to implement a highly simplified and barebones version of minecarts. Minecarts are an entity that travels along rails placed by the player. They can only move along rails, and have a velocity that can be affected by gravity, or other external forces. Minecarts are a perfect scenario to demonstrate Canary because they involve interaction between an entity and blocks in the world.

Our implementation of minecarts is highly simplistic, lacking most of the features of vanilla Minecraft. What we have implemented is that placing a minecart on top of a rail will cause it to travel in a straight line along the rail until it is no longer over a rail. As can be seen in Figure 5.1.1.1, we created 3 tests for our minecart implementation.



Figure 5.1.1.1: From left to right, test 1, test 2, and test 3 before being run

Test 1: straight line track

This test aimed to verify that when a minecart was placed on rails, it would travel along those rails.

The structure for the test consists of 3 rails on top of a layer of stone blocks. The code for this test is shown in Listing CD.1.

```
``java
// Listing CD.1
@InWorldTest
public void straightConstantSpeed(TestEnvironment env) {
    // spawns the minecart on top of the first rail (coordinates are relative to the structure origin)

    Pos starting_minecart_pos = new Pos(1.5, 1, 1.5);

    var minecart = env.spawnEntity(MinecartEntity::new, starting_minecart_pos);
    // expect the minecart to be 3 blocks over in the x direction
    env.expect(minecart).toBeAt(starting_minecart_pos.withX(x -> x + 3));
}

```

```

After running, test 1 passes. Shown in Figure 5.1.1.2 is the test after passing, with the green beacon to indicate success.




Figure 5.1.1.2: Test 1 after passing

### Test 2: No Track

This test verifies that the minecart only moves forward when it is on top of rails. The structure and code are nearly identical to test 1, except the structure associated with it does not have the rails for the minecart. For the test code, the only change is to the assertion, which now contains a ‘NOT’ (Listing CD.2).

```
```
``java
// Listing CD.2
env.expect(minecart).not().toBeAt(starting_minecart_pos.withX(x -> x + 2));
````
```

This test also passes, although it passes instantly, demonstrating a current shortcoming of our assertion engine. Because the only assertion contains a ‘NOT’, the first tick of the test results in a ‘SOFT\_PASS’, meaning the test stops executing. This issue could be addressed by allowing users to specify a minimum lifetime for the test, giving time to check for unwanted behavior.

### Test 3: Gravity

This tests for a behavior that is currently not implemented correctly: the ability to fall off one rail and land on another lower down. Currently, the minecart gets stuck at the top due to incorrect physics properties.




Figure 5.1.1.3: Test 3 after timing out

We see in Figure 5.1.1.3 that when the test fails, the beacon becomes red, and a book stand is placed. Right clicking on the book stand, as shown in Figure 5.1.1.4, gives the reason for test failure. In this case, the test timed out.




Figure 5.1.1.4: The error message for test 3

This demo feature implementation and testing shows the basic functionality of Canary. The process of creating the demo gave us the opportunity to use Canary like the intended users would. From this perspective, we found Canary to overall be very usable. Apart from some bugs and quirks as a result of the somewhat unpolished state of Canary, all features worked as expected. The test builder, assertion engine, and test executor do most of the hard work for users, letting them focus on high level thinking about how to define their tests.

## 5.2 Future Work

Canary implements a usable set of utilities to test many common features surrounding players, blocks, commands, and entities. However, creating a fully-featured testing framework for a complex environment like a Minecraft server is a non-trivial task. There are plenty of features which would improve the workflow, usability, and usefulness of Canary.

### 5.2.1 Unified workflow

The workflow in Canary involves creating a test structure inside the game and then switching over to a code editor to write the test code. In addition, running a test requires restarting the sandbox server or hot reloading the test code, which is notoriously problematic on the Java Virtual Machine (JVM). This context switch can cause a reduction in productivity and can be mitigated. The Mojang GameTest framework uses a combination of server commands and Behavior Packs (JavaScript) to execute tests. Using server commands in Minestom can be difficult given that no server commands are provided by default. As discussed, Canary introduces a number of utility commands for editing structures in the test builder; however, they do not extend to replacing test code acting on the structures. A more fluent solution would allow the user to “program” the tests inside the game using the command block feature of Minecraft. In this case, the tester would be able to edit the tests without ever leaving the game or worrying about reloading their Java code.

### 5.2.2 Testing Minestom Internally

Canary has a strong focus on testing applications built on top of Minestom, and it assumes that Minestom features are properly working. At the time of writing, however, Minestom does not have strong testing of internal components. Canary could help fill this gap. Testing internals are not dissimilar to testing applications built on top, and many principals would stay accurate. When testing in userland, you may generally ignore the client because the server acts as an adapter and is assumed to be working. To test the server, the client must be included. A system such as the attempt at creating a mock client for Minestom can test known behavior of the client. Unfortunately, the client is closed source and written by Mojang, meaning that changes to the client are not widely published, so features can break at any point. Alternatively, using a running client to test using the real behavior is both more accurate, and tests for regression when the client changes internally.

### 5.2.3 Behavior-driven Development

Behavior-driven development (BDD) is a practice similar to TDD where testing becomes an integral part of the writing process. Where BDD differs, however, is how tests are written; TDD takes a traditional approach of self-contained test functions, whereas BDD works in terms of "scenarios". These scenarios describe a series of events and expected behavior. A sample can be seen in Listing 5.2.3.1 for a store application where the scenarios describe what happens to an item when returned to a store. BDD provides two key benefits: tests which provide concrete usage examples, and collaboration with non-developers to create a shared understanding of functionality.

```
““
// Listing 5.2.3.1
// Adapted from https://en.wikipedia.org/wiki/Behavior-driven_development
Title: Returns go to inventory.
As a store owner,
I want to add items back to inventory when they are returned,
so that I can track inventory.
Scenario 1: Items returned for refund should be added to inventory.
Given that a customer previously bought a black sweater from me
And I have three black sweaters in inventory,
When they return the black sweater for a refund,
Then I should have four black sweaters in inventory.
““
```

In Canary, a user could create a test structure and then write a scenario for it in the game using a library of statements which have been defined by developers and Canary itself. This would solve the context switching issue discussed previously and allow the user to go from inception to execution all without leaving the game. Furthermore, since scenarios are designed to be abstract statements of expected functionality,

non-developers could contribute. This could lead to a workflow where a Canary sandbox is always running, and multiple people (developers or not) would join to test out functionality and create test cases while doing so, all working towards a shared understanding of the intended functionality.

## 6 Conclusion

Canary is a framework with surgical precision in scope. Because it is created specifically for Minestom, it can integrate extremely closely, providing users a seamless and easy to use experience. Taking inspiration from other projects with similar goals, we have created an end-to-end testing framework capable of becoming the standard for high level testing in Minestom. Indeed, the methods shown here, particularly the solution for making assertions on events happening in the future, could be applied to other projects inside or outside the game development space. While there is more to improve, Canary has opened up a discussion about testing in the Minestom ecosystem. It has shown the community that testing does not have to be difficult or time-consuming, and it will increase reliability of software considerably.

As of the writing of this paper, Canary has not had a full and complete release. The project will continue to be developed afterwards, addressing current issues as well as allowing for new methods of testing. We have shown it to be a simple framework for testing complex behavior, and we believe that testing can and will become a standard for Minestom projects in the future.

## 7 Appendix A: Java Classes & Runtime Modification

When a Java class is requested by another class, for example by calling its constructor, the class is loaded into the Java Virtual Machine (JVM) by a ‘`java.lang.ClassLoader`’. The language allows multiple class loaders to exist in a tree, which imposes restrictions on which classes can be accessed where. Class loaders form a bottom-up tree structure where class loaders know their parent, but not their children. Using the class loader setup in Figure AA.1 as an example, class loader A and B both have a shared parent of the root class loader. All classes must be loaded into a single class loader, which determines whether they are able to access other classes. For example, ‘`ClassA`’ is loaded into class loader A. Classes may only access other classes in their own class loader or in a parent. In this case, ‘`ClassA`’ may access ‘`Main`’ and ‘`StringUtil`’, but not ‘`ClassB`’. In the same respect, ‘`Main`’ may not access ‘`ClassA`’ or ‘`ClassB`’. When a class loader is instructed to load a class, it will first attempt to load the class from a parent class loader, and then try to load it in itself. In Listing AA.1, this means that if class loader A requests ‘`ClassA`’, it will return the loaded version. However, if the root class loader requests ‘`ClassA`’, a new, unique, instance will be loaded onto the root class loader. Having duplicate instances of the same class can cause significant problems because they are not convertible (for example, an instance of ‘`ClassA`’ from class loader A cannot be assigned to a variable of type ‘`ClassA`’ from the root class loader). This system is relevant to Minestom because Minestom uses the Mixin framework.




Figure AA.1: Example class loader hierarchy

Mixin is a framework developed by the SpongePowered team for runtime modification of Java classes. It hooks into the class loading process to modify the bytecode of the loading class. This allows a user to modify an existing class of a library without having to compile a new version or fork the project. Listing AA.2 shows an example Mixin which will add a system print to the first line of ‘CanaryTestEngine#discover‘ at runtime. Listing AA.3 shows the definition of ‘CanaryTestEngine#discover‘ at compile time and then at runtime after the Mixin has been applied. Minestom includes the Mixin framework so that extensions (Minestom’s plug-in system) are able to modify Minestom classes without requiring a fork of the project.

```

``java
// Listing AA.2

@Mixin(CanaryTestEngine.class)
public class CanaryTestEngineMixin {
 @Inject(method = "discover", at = @At("HEAD"))
 public void discover(EngineDiscoveryRequest discoveryRequest, UniqueId uniqueId, CallbackInfoReturnable<CanaryEngineDescriptor> cir) {
 System.out.println("Printed during discovery");
 }
}
``

``java
// Listing AA.3

@Override
public CanaryEngineDescriptor discover(EngineDiscoveryRequest discoveryRequest, UniqueId uniqueId)
{
 return CanaryDiscoverer.discover(discoveryRequest, uniqueId);
}
@Override
public CanaryEngineDescriptor discover(EngineDiscoveryRequest discoveryRequest, UniqueId uniqueId)

```

```

{
 System.out.println("Printed during discovery");

 return CanaryDiscoverer.discover(discoveryRequest, uniqueId);

}
```

```

The presence of this system adds a significant complexity to Canary. Mixin requires that all classes are loaded into a Mixin-enabled class loader so that it can intercept and modify the classes as they are loaded. To enable this, Minestom provides a bootstrap class to set up the Mixin environment and then invoke the developers' main class. Listing AA.4 shows an example initialization for a Minestom project. Unfortunately, this system requires that the developer is in charge of where classes are loaded. Inside a JUnit test engine, this is not the case. JUnit instantiates the test engine class and manages loading the discovered test classes into its own class loader. This means that Mixins will not be applied, and user features using them will likely not work. To get around this problem, Canary keeps two isolated class loaders: the JUnit class loader where the test engine is loaded and working, and the Minestom server class loader where the server and tests are running. Since test classes are loaded into the JUnit class loader but executed in the Minestom class loader, they must be reloaded before execution. This system adds significant complexity and safety requirements to Canary (see Appendix B); however, it comes with an advantage past enabling Mixin support. Because the test engine is independent from the server, it is possible to start multiple independent Minestom servers which could provide complete isolation between running tests.

```

``java
// Listing AA.4

public class BootstrapMain {
    public static void main(String[] args) {
        // com.example.Main.main(String[]) will be invoked after being loaded into the Mixin class loader.
        Bootstrap.bootstrap("com.example.Main", args);
    }
}
```

```

This project and report were written for Minestom targeting Minecraft version 1.17. The game is now on 1.18, and during the 1.18 update, Minestom removed Mixin support. This means that the system will no longer be required for Minestom versions going forwards unless Canary itself needs to use Mixins to modify Minestom for any reason.

## 8 Appendix B: Tooling

Canary has some complex systems to manage class loaders and assertions, but to aid in development and usage, we created some external supporting tooling. A Gradle plugin and a Java compiler plugin were

both created to help simplify usage and reduce errors. The Gradle plugin is mainly for external users, whereas the compiler plugin is intended for internal use only.

The Canary Gradle plugin is a simple plugin to facilitate development when using Canary in an external project. The obvious task of the plugin is to set up Canary as a test dependency. Canary should never be added to the main class path, so it is only set up on the test class path. Additionally, in the future the API will likely be split away from the JUnit test engine implementation, creating a more complex dependency setup. Next, the plugin creates a task to start the sandbox server. The sandbox server may be started manually by invoking its main class (`com.mattworzala.canary.internal.launch.SandboxLauncher`); however, there are some required command line arguments which would have to be added manually. These options are configurable through the Gradle plugin. An example usage can be seen in Listing AB.1. The plugin is added via its ID (`com.mattworzala.canary`) and then may be configured through the `canary` block. In this sample, only the Canary version is being configured.

```
```
“‘kotlin
// Listing AB.1
// build.gradle.kts
plugins {
    id("com.mattworzala.canary")
}
tasks {
    canary {
        version = "-SNAPSHOT"
    }
}
```
``
```

While the Gradle plugin facilitates external usage, there is some room for internal tooling as well. As described in Appendix A, Canary requires strict isolation between the JUnit Classloader and the Minestom Classloader. As such, a reference to a JUnit class, such as `CanaryTestEngine` (the JUnit engine implementation) is invalid from a Minestom server class, such as `HeadlessServer`. To prevent accidental references, we created a Java compiler plugin to enforce this isolation. A Java compiler plugin is similar to an annotation processor; however, it is given direct access to the Abstract Syntax Tree (AST) and may manipulate it with much more freedom. The downside, however, is that the `javac` Application Programming Interface (API) does not have the same backwards compatibility guarantees as the annotation processor API. In this case, a compiler plugin is required because annotation processors do not have access to import statements. The Canary compiler plugin (or safety checks) works by looking at import statements and the `@Env` annotation. The annotation is used on class definitions and tells the safety checker what environment (class loader) the class may be used in. The valid options are `GLOBAL` for classes available in both class loaders and `PLATFORM / MINESTOM` for JUnit/Minestom respectively. With this information, errors can be generated from invalid import statements. Figure AB.2 shows an error generated by the safety checker when importing `CanaryTestEngine` from `HeadlessServer`.

```
““
// Figure AB.2
HeadlessServer.java:6: error: Cannot import com...CanaryTestEngine (PLATFORM) from MINESTOM
import com.mattworzala.canary.internal.junit.CanaryTestEngine;
^
““
```